

From Desktop to Exascale: All in on GPUs



Michael A. Heroux
Senior Scientist, Sandia National Laboratories
Former Director of Software Technology, ECP
PI, PESO, Post-ECP Software-Ecosystem Stewardship Project
SIAM Fellow

SIAM Supercomputing Spotlights Series
September 11, 2024



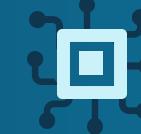
Exascale Computing Project



Maintain
international
leadership
in HPC



Promote the
health of the
US HPC industry



Deliver a **sustainable**
software ecosystem
used and maintained
for years to come



Ensure that exascale
systems can be used
to deliver **mission-**
critical applications



7-year, \$1.8B

US Department of Energy
project funded 1000+ people
at national labs, universities,
US industries

This research was supported by the Exascale
Computing Project (17-SC-20-SC), a collaborative effort
of the U.S. Department of Energy Office of Science and
the National Nuclear Security Administration.



Application Development

- Develop and enhance the predictive capability of applications,
25 applications, 6 Co-Design Centers

Software Technology

- Deliver expanded and vertically integrated software stack, **70 unique products**

Hardware and Integration

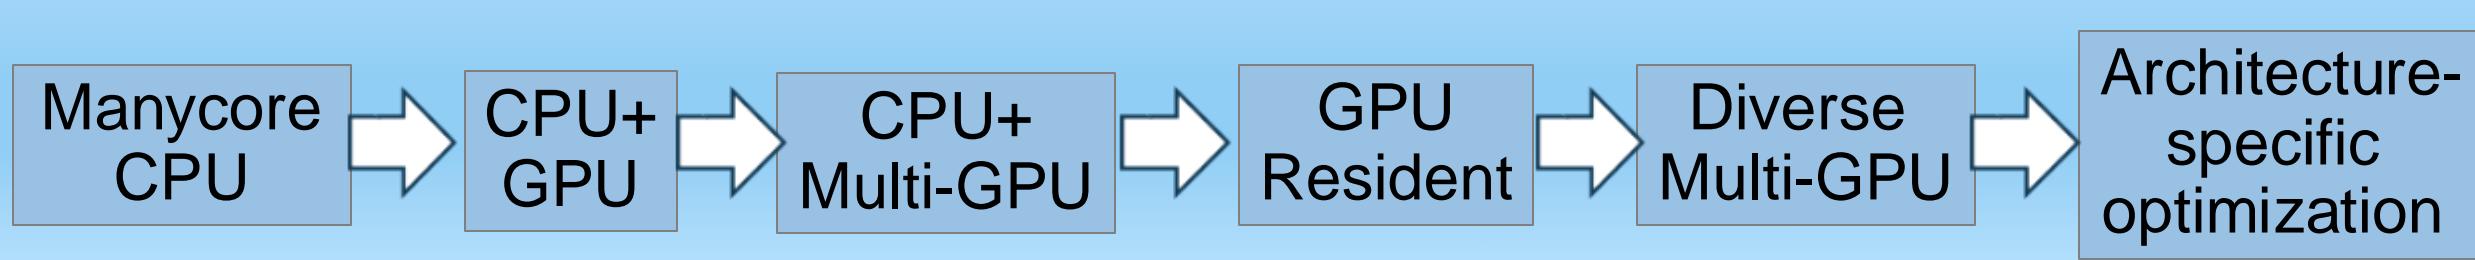
- Application integration and software deployment to facilities, exascale node
and system design, **6 US HPC vendors**

ECP Application Development

National security	Energy security	Economic security	Scientific discovery	Earth system	Health care
Next-generation, stockpile stewardship codes Reentry-vehicle-environment simulation	Turbine wind plant efficiency Design and commercialization of SMRs	Additive manufacturing of qualifiable metal parts Reliable and efficient planning	Cosmological probe of the standard model of particle physics Validate fundamental laws of nature	Accurate regional impact assessments in Earth system models Stress-resistant crop analysis and catalytic	Accelerate and translate cancer research (partnership with NIH)
					

The 24 AD application projects

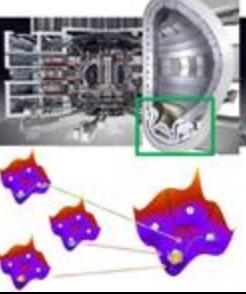
- Include **62 separate codes**
- Represent over **10 million lines of code**
- In some cases support large user communities
- Mostly started with MPI or MPI+OpenMP on CPUs



ECP libraries and tools supported diverse applications across multiple architectures

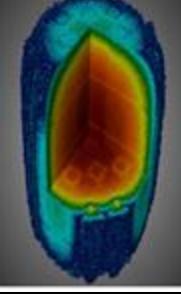
Project/PI	EXAALT: Molecular Dynamics Danny Perez
Challenge Problem	Damaged surface of Tungsten in conditions relevant to plasma facing materials in fusion reactors <ul style="list-style-type: none"> • 100,000 atoms • T=120K
FOM Speedup	398.5
Nodes Used	7000

ST/CD Tools Used in KPP Demo: Kokkos, CoPA



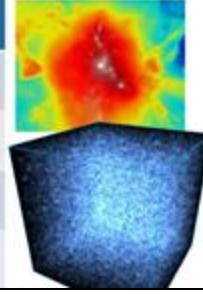
Project/PI	ExaSMR: Small Modular Reactors Steve Hamilton
Challenge Problem	NuScale-style Small Modular Reactor (SMR) with depleted fuel and natural circulation <ul style="list-style-type: none"> • 213,860 Monte Carlo tally cells/6 reactions • 5.12×10^{12} particle histories/cycle, 40 cycles • 1098×10^6 CFD spatial elements • 376×10^7 CFD degrees of freedom • 1500 CFD timesteps
FOM Speedup	70
Nodes Used	6400

ST/CD Tools Used in KPP Demo: CEED
Additional: Trilinos



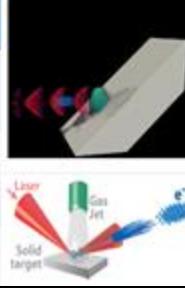
Project/PI	ExaSky: Cosmology Salman Habib
Challenge Problem	Two large cosmology simulations <ul style="list-style-type: none"> • gravity-only • hydrodynamics
FOM Speedup	271.65
Nodes Used	8192

ST/CD Tools Used in KPP demo: none
Additional: CoPA, VTK-m, CINEMA, HDF5.0



Project/PI	WarpX: Plasma Wakefield Accelerators Jean-Luc Vay
Challenge Problem	Wakefield plasma accelerator with a 1PW laser drive <ul style="list-style-type: none"> • 6.9×10^{12} grid cells • 1.4×10^{12} macroparticles • 1000 timesteps/1 stage
FOM Speedup	500
Nodes Used	8576

ST/CD Tools Used in KPP Demo: AMReX, libEnsemble
Additional: ADIOS, HDF5, VTK-m, ALPINE




ECP Software Technologies

Prepare SW stack for scalability with massive on-node parallelism

Extend existing capabilities when possible, develop new when not

Guide, and complement, and integrate with vendor efforts

Develop and deliver high-quality and robust software products

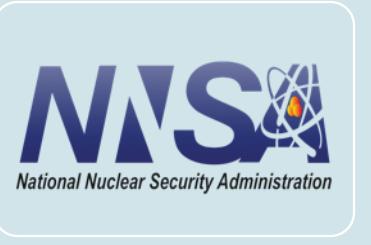
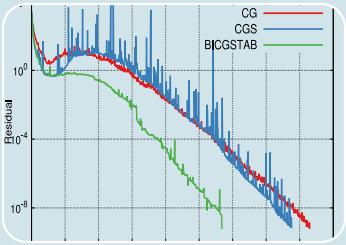
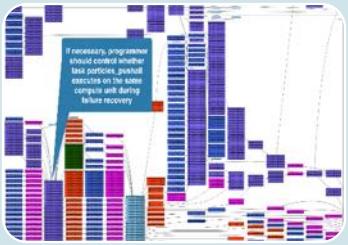
70 software products across 6 technical areas



ECP Impact – Portable Libraries and Tools for Accelerators



ECP ST six technical areas



Programming Models & Runtimes

- Enhance and get ready for exascale the widely used MPI and OpenMP programming models (hybrid programming models, deep memory copies)
- Development of performance portability tools (e.g. Kokkos and Raja)
- Support alternate models for potential benefits and risk mitigation: PGAS (UPC++/GASNet), task-based models (Legion, PARSEC)
- Libraries for deep memory hierarchy and power management

Development Tools

- Continued, multifaceted capabilities in portable, open-source LLVM compiler ecosystem to support expected ECP architectures, including support for F18
- Performance analysis tools that accommodate new architectures, programming models, e.g., PAPI, Tau

Math Libraries

- Linear algebra, iterative linear solvers, direct linear solvers, integrators and nonlinear solvers, optimization, FFTs, etc
- Performance on new node architectures; extreme strong scalability
- Advanced algorithms for multi-physics, multiscale simulation and outer-loop analysis
- Increasing quality, interoperability, complementarity of math libraries

Data and Visualization

- I/O via the HDF5 API
- Insightful, memory-efficient in-situ visualization and analysis – Data reduction via scientific data compression
- Checkpoint restart

Software Ecosystem

- Develop features in Spack necessary to support all ST products in E4S, and the AD projects that adopt it
- Development of Spack stacks for reproducible turnkey deployment of large collections of software
- Optimization and interoperability of containers on HPC systems
- Regular E4S releases of the ST software stack and SDKs with regular integration of new ST products

NNSA ST

- Open source NNSA Software projects
- Projects that have both mission role and open science role
- Major technical areas: New programming abstractions, math libraries, data and viz libraries
- Cover most ST technology areas
- Subject to the same planning, reporting and review processes

ECP-sponsored Software Products: A Sample

Programming Models & Runtimes

GPU-specific kernels

- Isolate the computationally-intensive parts of the code into CUDA/HIP/SYCL kernels.
- Refactoring the code to work well with the GPU is the majority of effort.

Loop pragma models

- Offload loops to GPU with OpenMP or OpenACC.
- Most common portability strategy for Fortran codes.

C++ abstractions

- Fully abstract loop execution and data management using advanced C++ features.
- Kokkos and RAJA developed by NNSA in response to increasing hardware diversity.

Four Parallel Node Programming Approaches

Co-design frameworks

- Design application with a specific motif to use common software components
- Depend on co-design code (e.g. CEED, AMReX) to implement key functions on GPU.

Efficiently utilizing GPUs goes far beyond typical code porting

Port Code

- Rewrite, profile, and optimize
- Memory coalescing
- Loop ordering
- Kernel flattening

Adapt Numerics

- Reduced synchronization
- Reduced precision
- Communication avoiding

Adapt Models

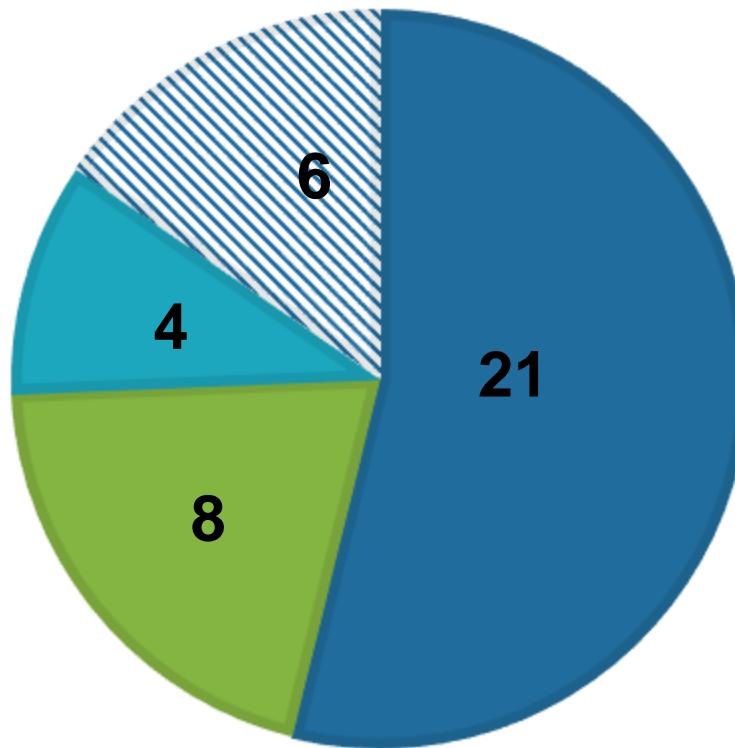
- Mathematical representation
- “On the fly” recomputing vs. lookup tables
- Prioritization of new physical models

ECP efforts have de-risked (but not removed) these activities for others

DOE Exascale Computing Project

ECP SOFTWARE PROJECTS USING C++: WITH CRITICAL DEPENDENCY ON

■ Kokkos Based ■ AMReX Based ■ RAJA Based ■ Other C++ Codes



- 17 Codes list Fortran as critical
- Some of those actually moved on
 - Not all of those run on GPUs
 - CUDA Fortran, OpenMP, OpenACC

Large Majority of C++ Codes Chose Abstraction over Vendor Models!

Kokkos Programming Model

What is Kokkos?

A C++ Programming Model for Performance Portability

- Implemented as a template library on top of CUDA, OpenMP, HIP, SYCL, ...
- Aims to be descriptive not prescriptive
- Aligns with developments in the C++ standard
- Replaces usage of CUDA, OpenMP, HIP, etc.

Expanding solution for common needs of modern science/engineering codes

- Math libraries based on Kokkos
- Tools which enable insight into Kokkos

It is Open Source

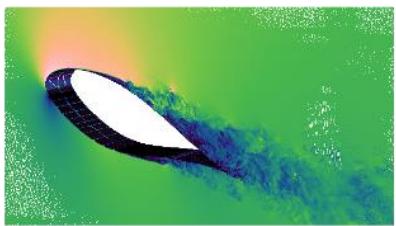
- Maintained and developed at <https://github.com/kokkos>

It has many users at wide range of institutions.

Kokkos is NOT just for GPUs!

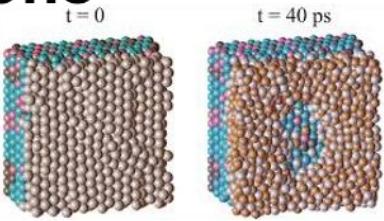


Applications

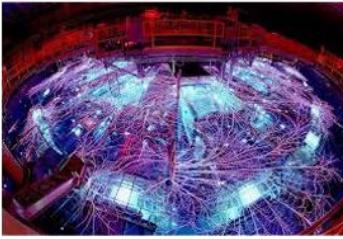


Computational
Fluid Dynamics

Libraries

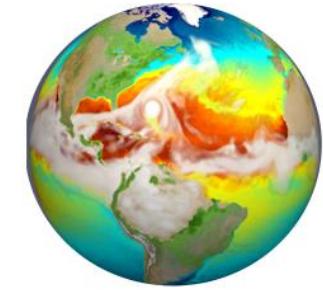


Molecular Dynamics



Electro Magnetics

Frameworks

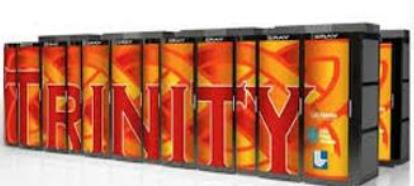


Climate Simulation

Kokkos



ORNL Frontier
AMD GPUs



LANL/SNL Trinity
Intel CPUs



ANL Aurora
Intel GPUs

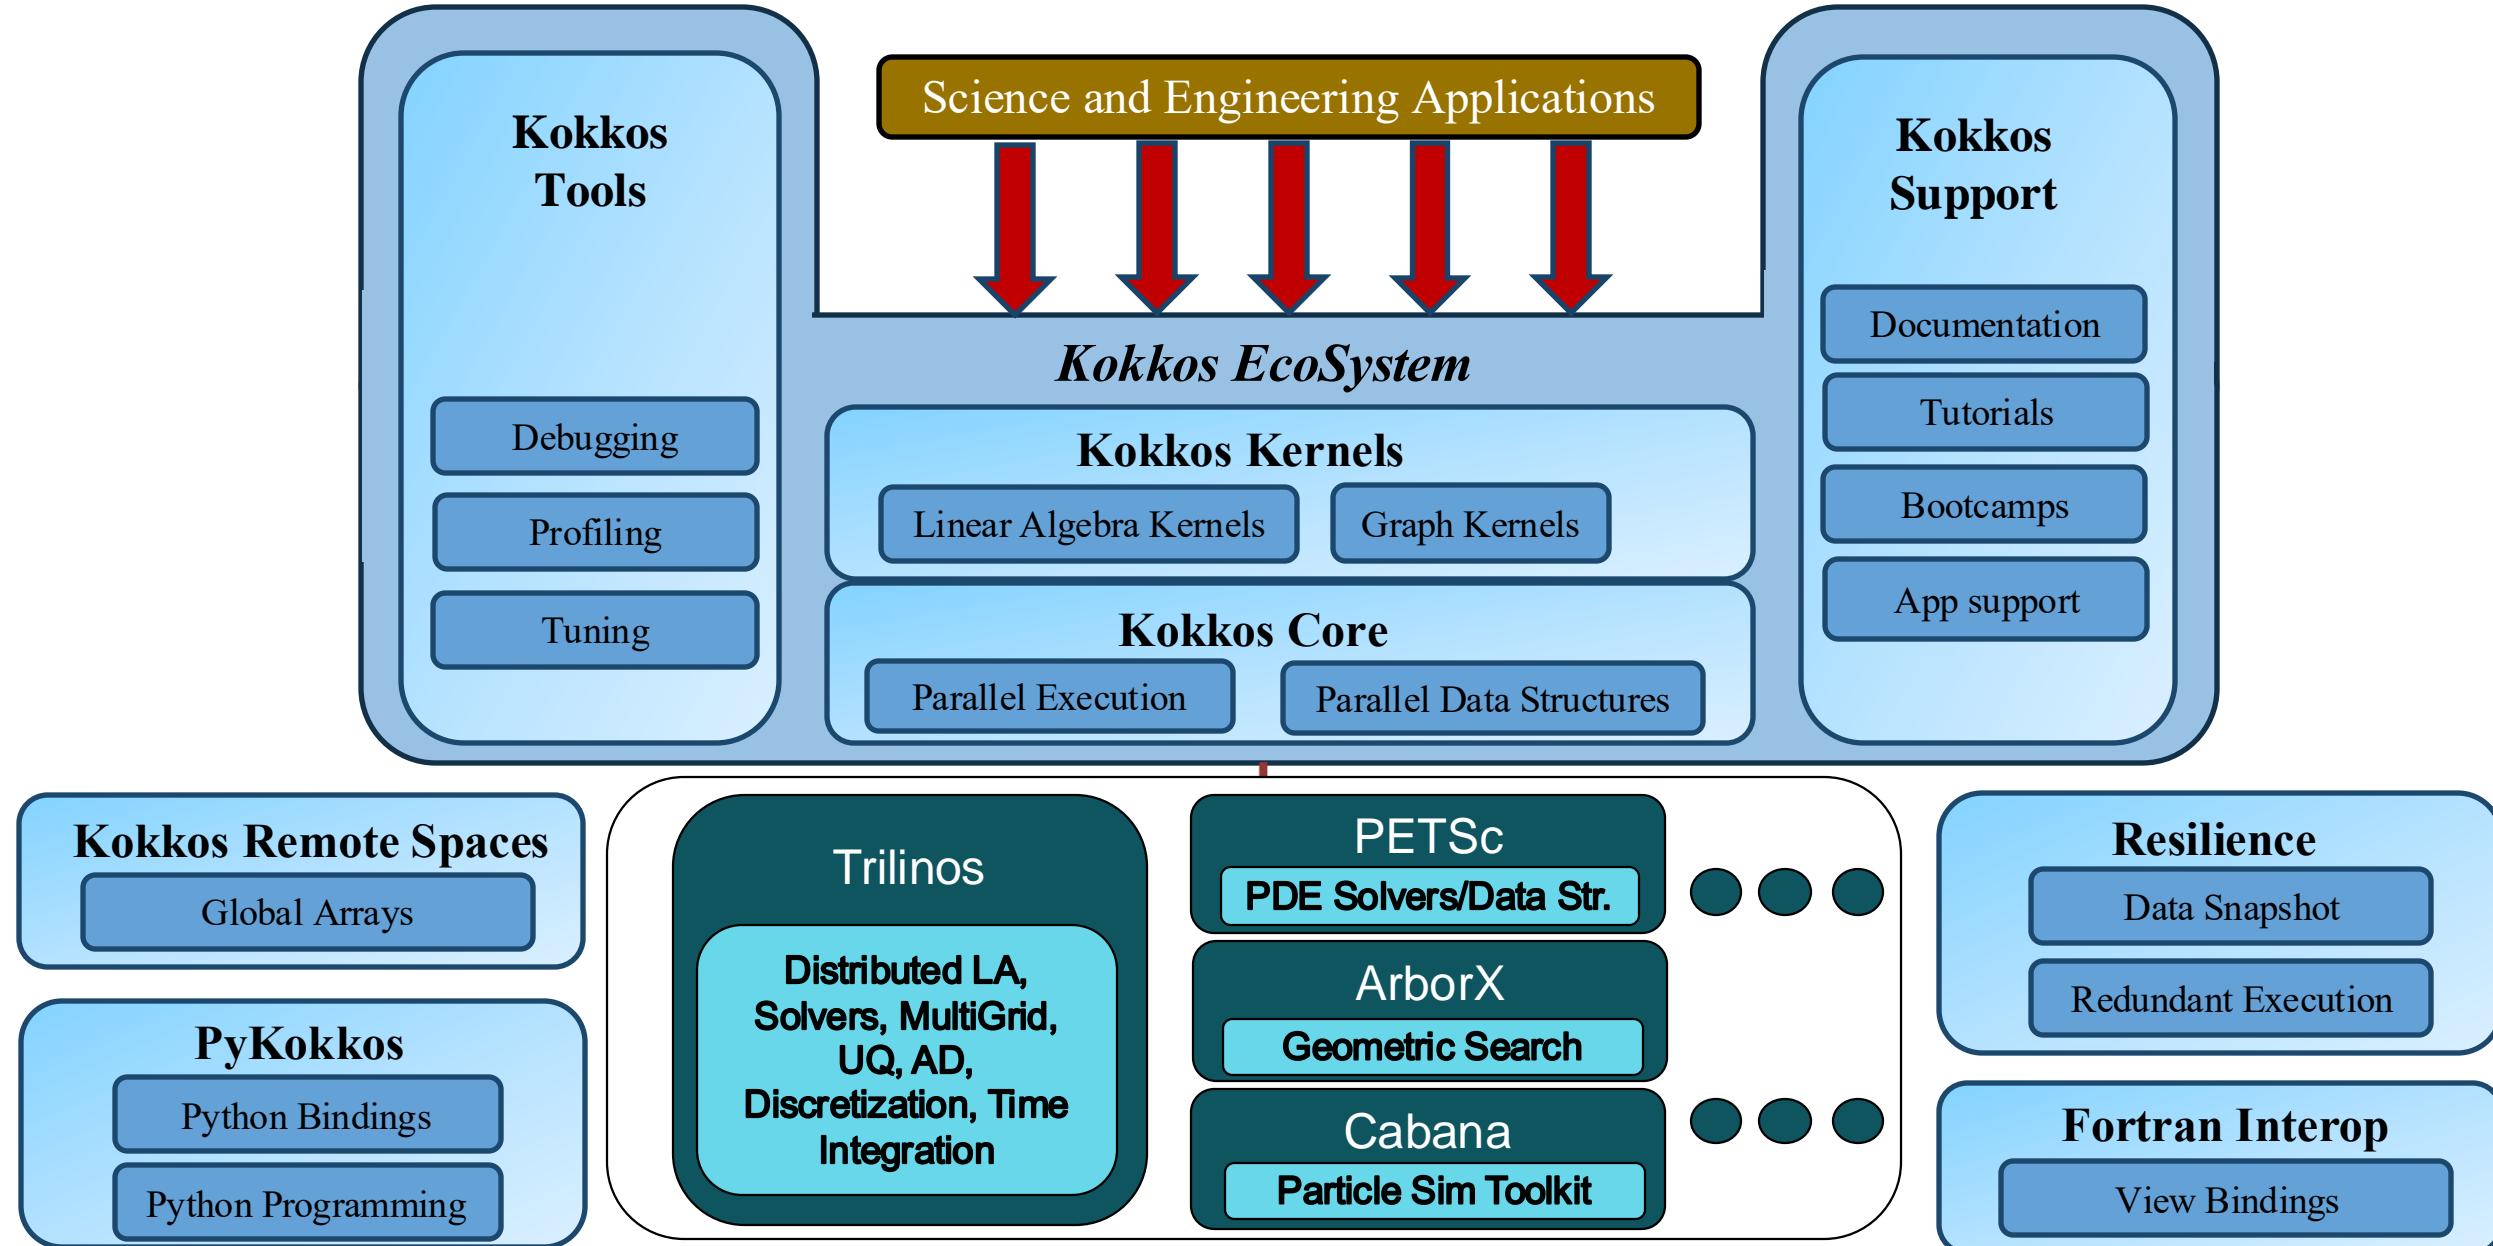


Riken Fugaku
ARM CPUs

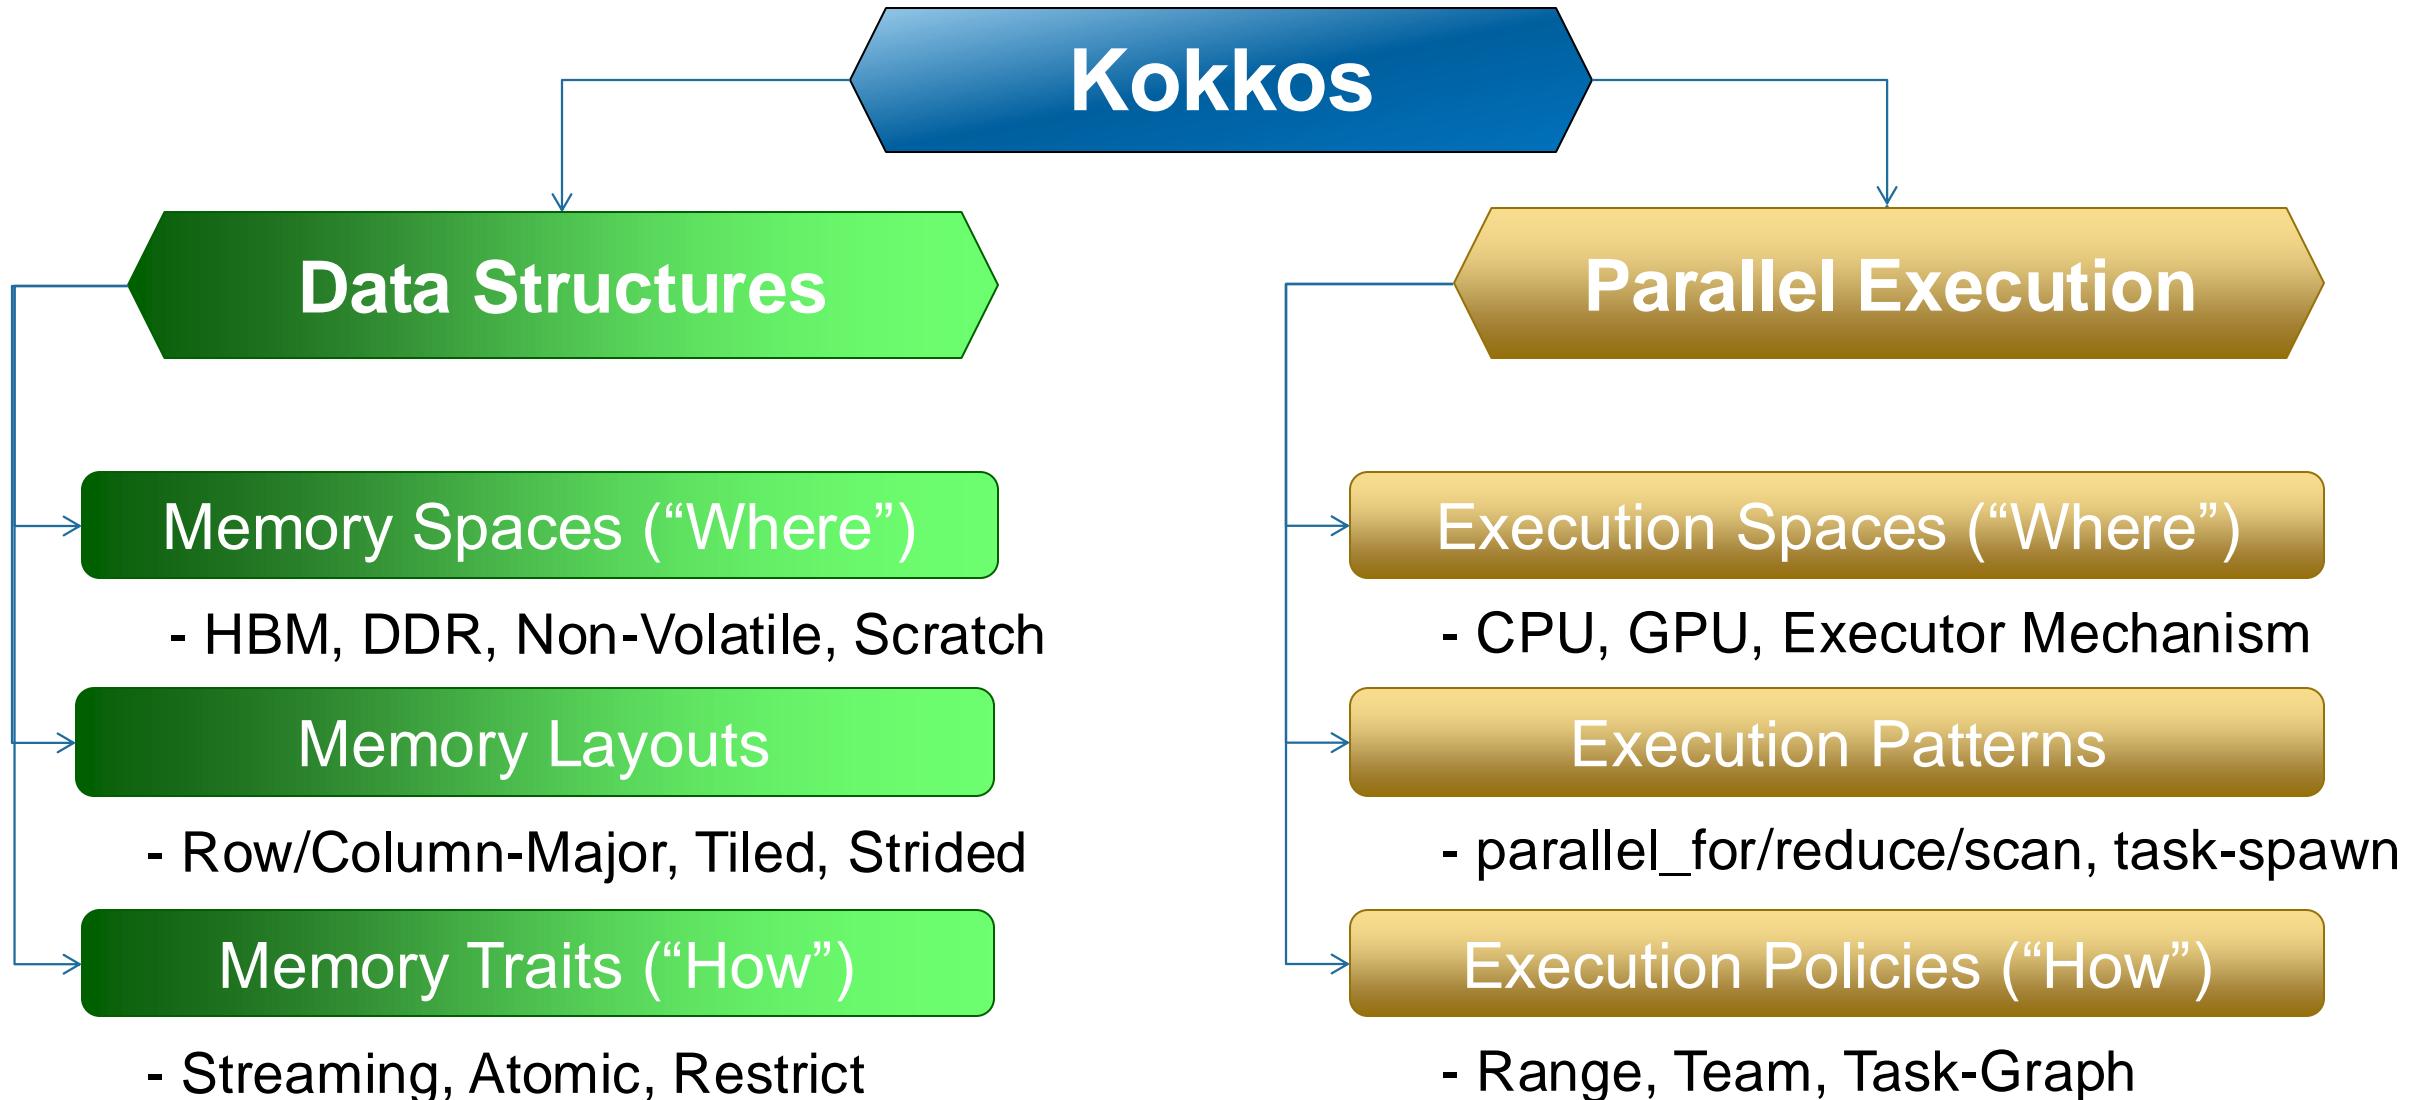


NERSC Perlmutter
NVIDIA GPUS

The Kokkos Ecosystem



Kokkos Core Abstractions



CG Solve: The AXPBY

Simple data parallel loop: Kokkos::parallel_for

Easy to express in most programming models

Bandwidth bound

Serial Implementation:

```
void axpby(int n, double* z, double alpha, const double* x,
           double beta, const double* y) {
    for(int i=0; i<n; i++)
        z[i] = alpha*x[i] + beta*y[i];
}
```

Parallel Pattern: for loop

Kokkos Implementation:

```
void axpby(int n, View<double*> z, double alpha, View<const double*> x,
           double beta, View<const double*> y) {
    parallel_for("AXpBY", n, KOKKOS_LAMBDA (const int i) {
        z(i) = alpha*x(i) + beta*y(i);
    });
}
```

String Label: Profiling/Debugging

Execution Policy: do n iterations

Loop Body

Iteration handle: integer index

Kokkos Support

The Kokkos Lectures

- 8 lectures covering most aspects of Kokkos
- 15 hours of recordings
- > 500 slides
- >20 exercises

Extensive Wiki

- API Reference
- Programming Guide

Slack as primary direct support

<https://kokkos.link/the-lectures>

- Module 1: Introduction
 - Introduction, Basic Parallelism, Build System
- Module 2: Views and Spaces
 - Execution and Memory Spaces, Data Layout
- Module 3: Data Structures and MDRangePolicy
 - Tightly Nested Loops, Subviews, ScatterView,...
- Module 4: Hierarchical Parallelism
 - Nested Parallelism, Scratch Pads, Unique Token
- Module 5: Advanced Optimizations
 - Streams, Tasking and SIMD
- Module 6: Language Interoperability
 - Fortran, Python, MPI and PGAS
- Module 7: Tools
 - Profiling, Tuning , Debugging, Static Analysis
- Module 8: Kokkos Kernels
 - Dense LA, Sparse LA, Solvers, Graph Kernels

Kokkos Kernels

BLAS, Sparse and Graph Kernels on top of Kokkos and its View abstraction

- Scalar type agnostic, e.g. works for any types with math operators
- Layout and Memory Space aware

Can call vendor libraries when available

Views contain size and stride information => Interface is simpler

// BLAS

```
int M, N, K, LDA, LDB; double alpha, beta; double *A, *B, *C;  
dgemm('N','N', M, N, K, alpha, A, LDA, B, LDB, beta, C, LDC);
```

vs.

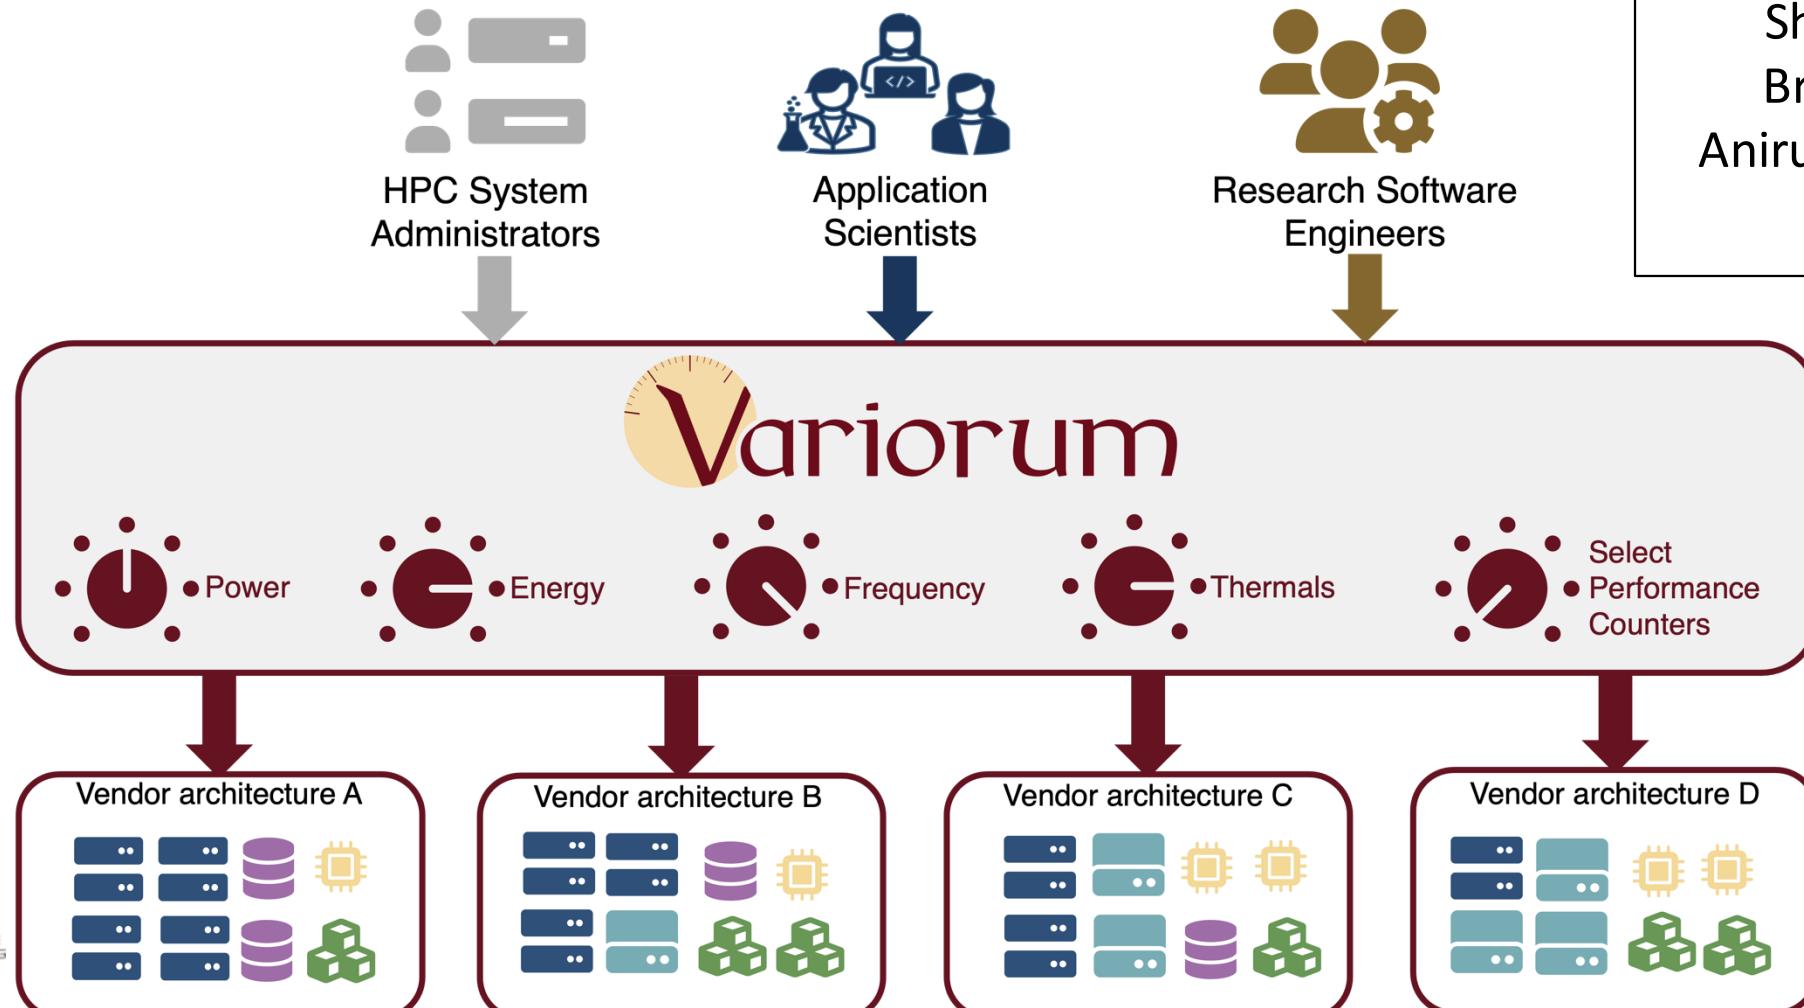
// Kokkos Kernels

```
double alpha, beta; View<double**> A,B,C;  
gemm('N','N', alpha, A, B, beta, C);
```

Development Tools

Variorum provides safe, user-space, vendor neutral access for all users: administrators, application scientists and RSEs

PI: Tapasya Patki
Team: Kathleen Shoga, Stephanie Brink, Eric Green, Aniruddha Marathe, Barry Rountree

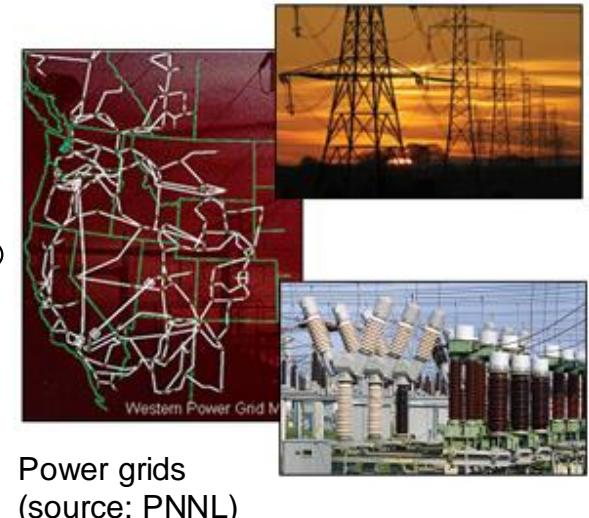
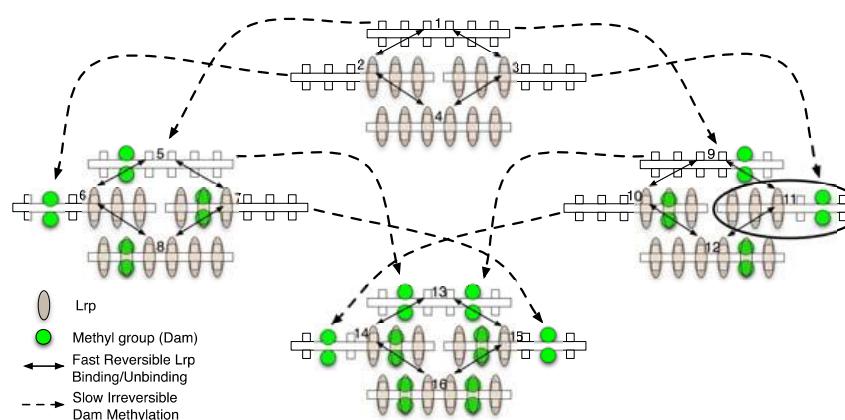
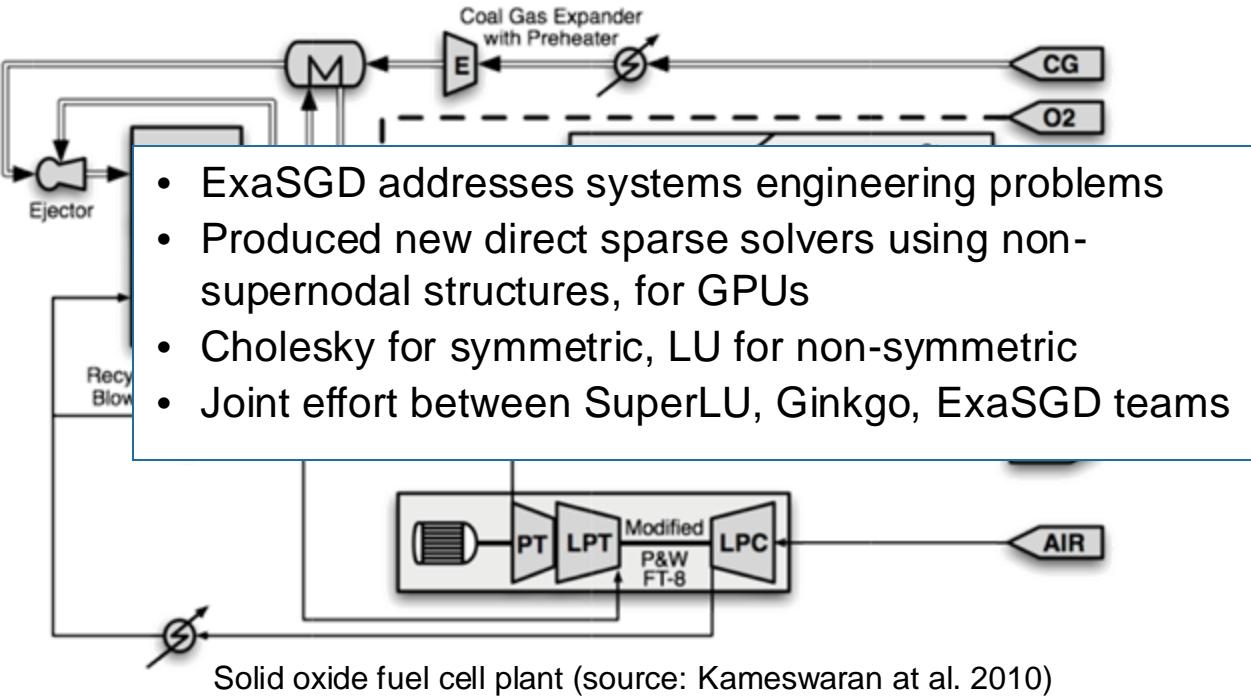
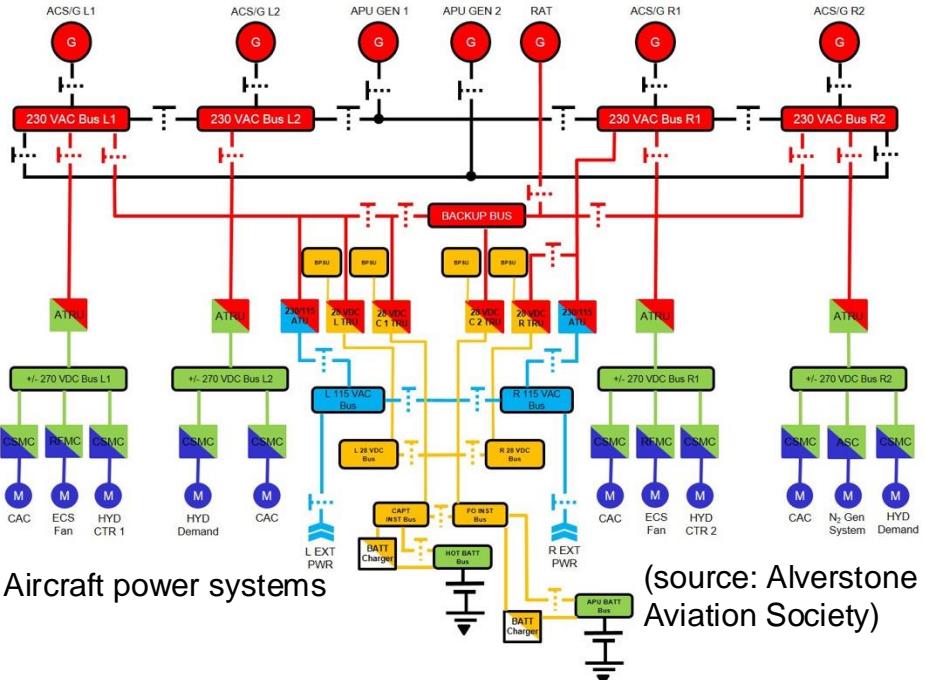


Variorum: Vendor-neutral user space library for power management

- Power management capabilities (and their interfaces, domains, latency, capabilities) widely differ from one vendor to the next, needing common interfaces
- Variorum: Platform-agnostic vendor-neutral, simple front-facing APIs
 - Evolved from *libmsr*, and designed to target several platforms and architectures
 - Abstract away tedious and chaotic details of low-level knobs
 - Implemented in C, with function pointers to specific target architecture
 - Integration with higher-level power management software through JSON
- Integrated with Flux, GEOPM, LDMS, Kokkos, Caliper and PowerAPI to enable a PowerStack
- Supported on all upcoming exascale systems (Aurora, Frontier, El Capitan) and several other supercomputers: architecture support includes CPU support for ARM, AMD, Intel, IBM; and GPU support for NVIDIA, AMD and Intel.

Math Libraries

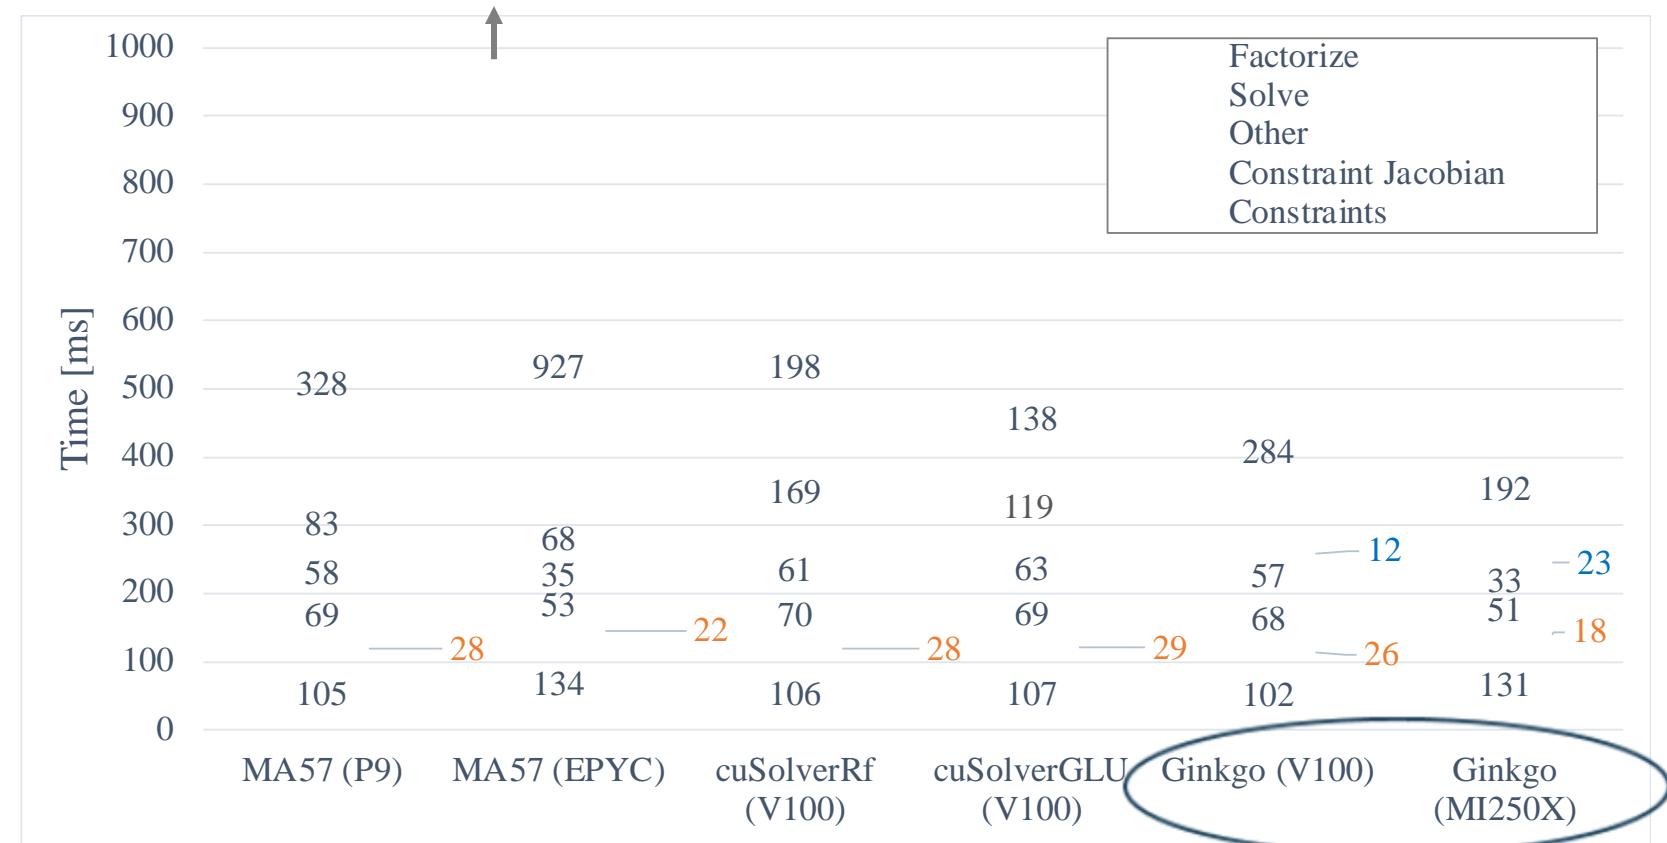
Systems Engineering Domain



Linear Solver Performance within Optimization Algorithm

Average per iteration times (including first iteration on CPU)

- Each GPU solution outperforms all CPU baselines
- Ginkgo performance improves on a better GPU
- Iterative refinement configuration affects linear solver performance and optimization solver convergence

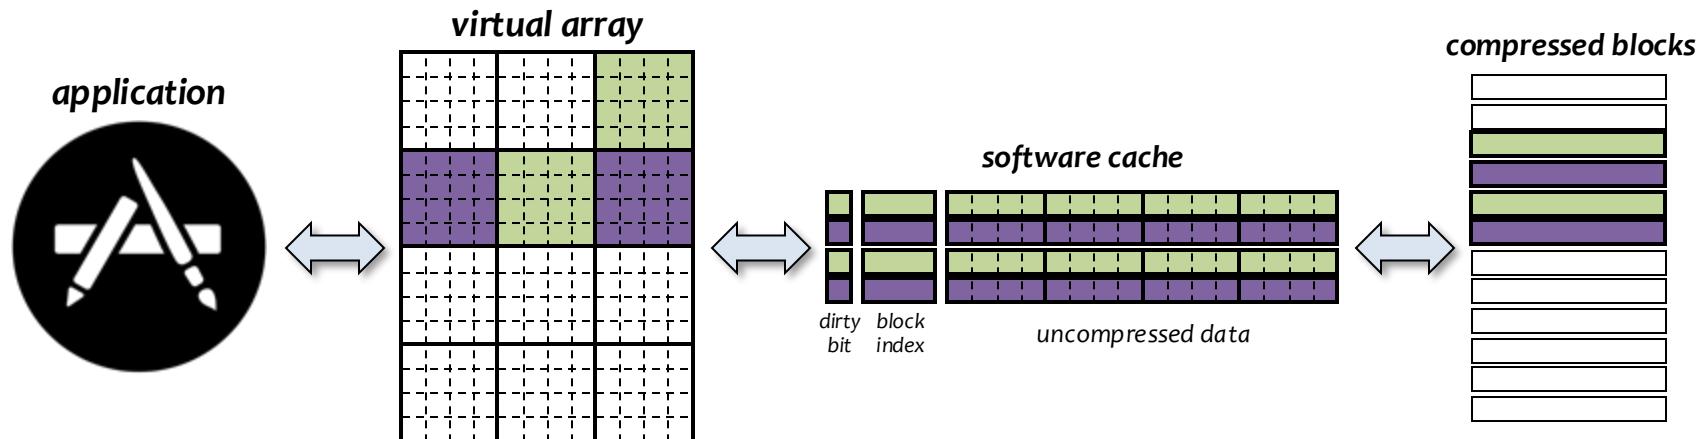


Ginkgo provides the first portable GPU-resident sparse direct linear solver for non-supernodal systems

Data and Visualization

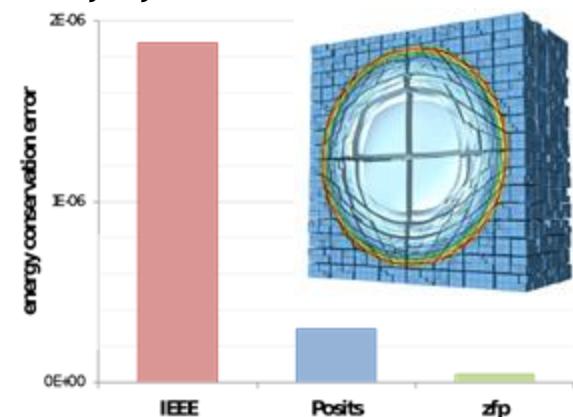
ZFP compressed multidimensional array primitive

Addressing growing gap of ops vs bw vs memory



- Fixed-length compressed blocks enable fine-grained read & write **random access**
 - C++ compressed-array classes hide complexity of compression & caching from user
 - User specifies per-array storage footprint in bits/value
- Absolute and relative **error tolerances** supported for offline storage, sequential access
- Fast, hardware friendly, and parallelizable: **150 GB/s throughput** on NVIDIA Volta
- **HPC tool support:** ADIOS, CEED, CONDUIT, HDF, intel, VTK

Inline compression boosts solution accuracy by 40x over IEEE in CEED code



And Many More...

- ECP generated a
 - Collection of portable GPU-capable libraries and tools for AMD, Intel, and NVIDIA devices
 - Designed for future adaptation to next-generation highly-concurrent node architectures
 - Foundation for others who will make the transition from CPU to GPU and beyond

ECP Software Technology works on products that apps need now and in the future

Key themes:

- Focus: GPU node architectures and advanced memory & storage technologies
- Create: New high-concurrency, latency tolerant algorithms
- Develop: New portable (Nvidia, Intel, AMD GPUs) software product
- Enable: Access and use via standard APIs

Legacy: A stack that enables performance portable application development on leadership platforms

Software categories:

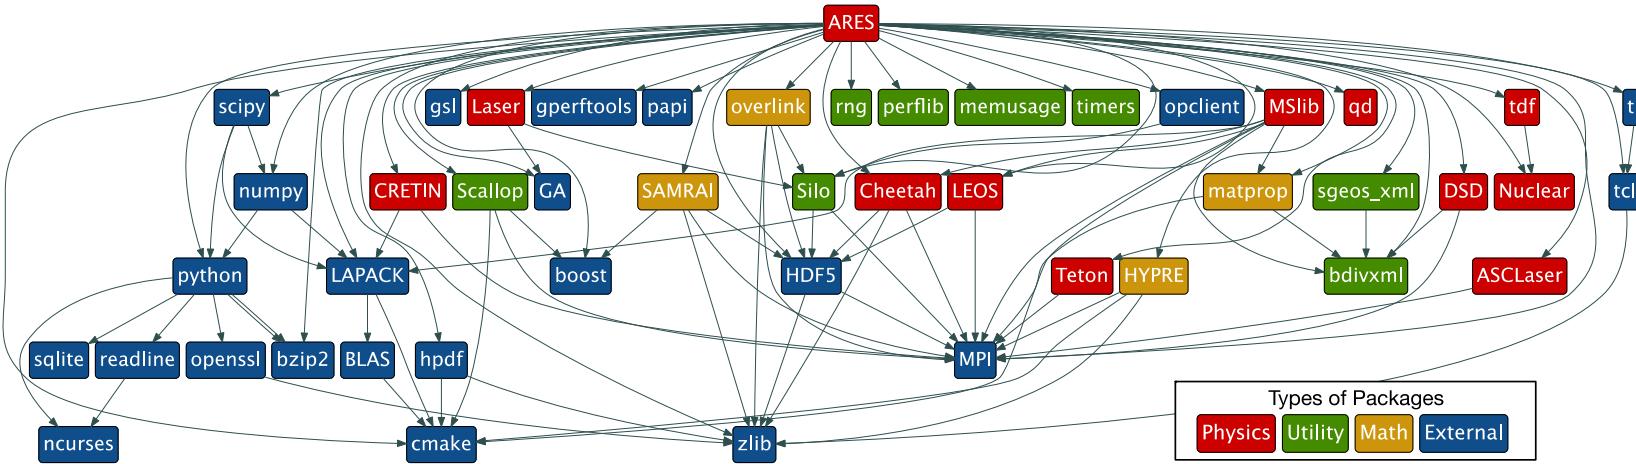
- **Next generation established products:** Widely used HPC products (e.g., MPICH, OpenMPI, PETSc)
- **Robust emerging products:** Address key new requirements (e.g., Kokkos, RAJA, Ginkgo, Spack)
- **New products:** Enable exploration of emerging HPC requirements (e.g., Variorum, zfp)



Example Products	Engagement
MPI – Backbone of HPC apps	Explore/develop MPICH and OpenMPI new features & standards
OpenMP/OpenACC –On-node parallelism	Explore/develop new features and standards
C++ Performance Portability Abstractions	Lightweight APIs for compile-time polymorphisms
LLVM/Vendor compilers	Injecting HPC features, testing/feedback to vendors
Perf Tools - PAPI, TAU, HPCToolkit	Explore/develop new features
Math Libraries: BLAS, sparse solvers, etc.	Scalable algorithms and software, critical enabling technologies
IO: HDF5, MPI-IO, ADIOS	Standard and next-gen IO, leveraging non-volatile storage
Viz/Data Analysis	ParaView-related product development, node concurrency

Spack: How we build and test our software

Spack automates the build of LLNL multi-physics codes



- ARES is a 1, 2, and 3-D radiation hydrodynamics code
- The ARES configuration shown here has **47 dependencies**
- ARES team runs nightly builds for 36 different configurations
 - 4 code versions:
 - **(C)urrent Production**
 - **(P)revious Production**
 - **(L)ite**
 - **(D)evelopment**

	Linux			BG/Q	Cray XE6
	MVAPICH	MVAPICH2	OpenMPI	BG/Q MPI	Cray MPI
GCC	C P L D			C P L D	
Intel 14	C P L D				
Intel 15	C P L D	D			
PGI		D	C P L D		C L D
Clang	C P L D			C L D	
XL				C P L D	

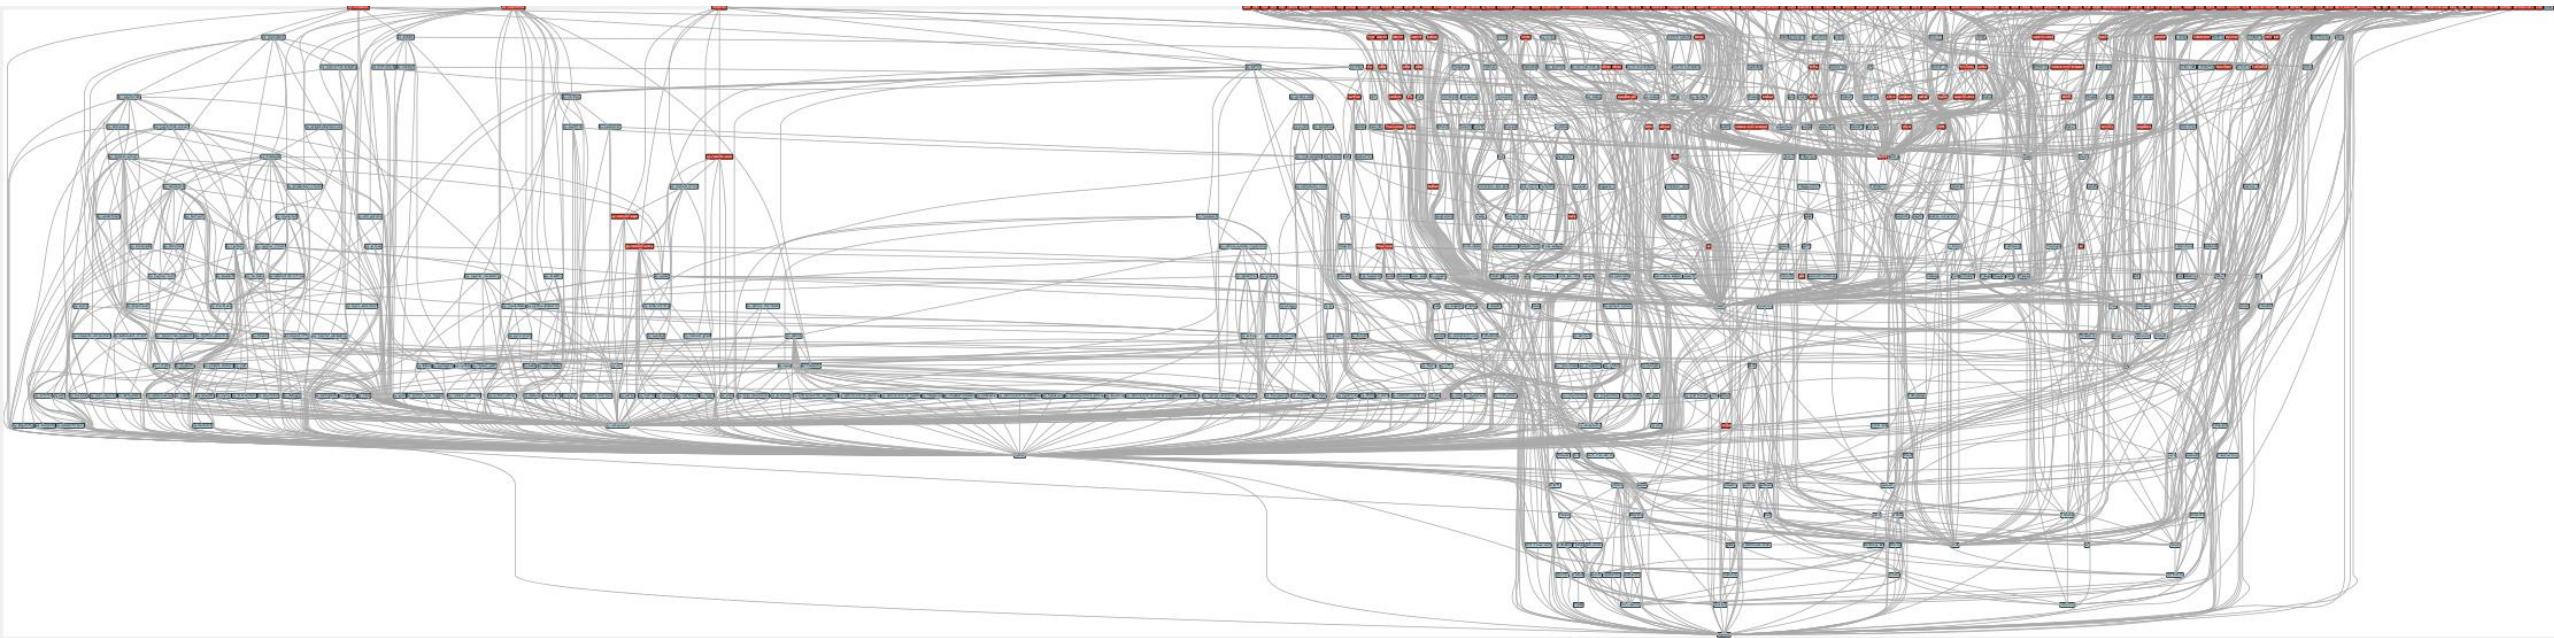
Spack enabled testing on 3 platforms, 6 compilers, and 3 MPI implementations.

Spack has been ambitious about customizability from the start: “Just say what you want to install”

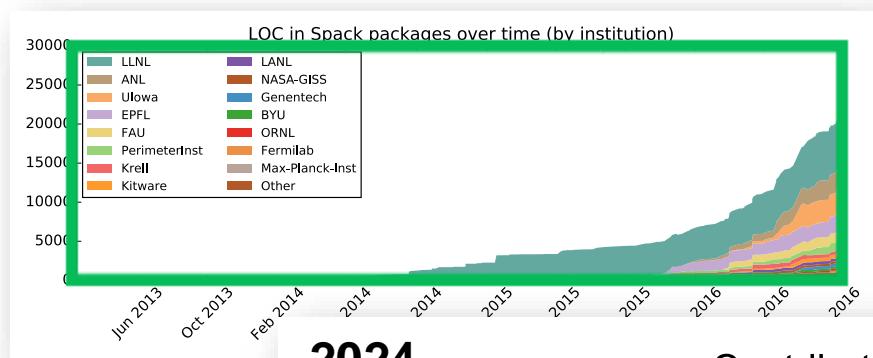
```
$ spack install mpileaks           unconstrained  
$ spack install mpileaks@3.3       @ custom version  
$ spack install mpileaks@3.3 %gcc@4.7.3    % custom compiler  
$ spack install mpileaks@3.3 %gcc@4.7.3 +threads  +/- build option  
$ spack install mpileaks@3.3 cppflags="-O3 -g3"  set compiler flags  
$ spack install mpileaks@3.3 target=cascadelake  set target microarchitecture  
$ spack install mpileaks@3.3 ^mpich@3.2 %gcc@4.9.3  ^ dependency constraints
```

- DOE people never settled for the “off-the-shelf” configuration
 - Always want to customize library versions, build options, flags
- Composition is a fundamental part of Spack’s model
 - Allows you to **swap components** like compilers, libraries, GPU runtimes, etc.
- Most distributions build a much more static stack

Now, the E4S stack is more than 600 packages!



- Red boxes are the packages in it (about 100)
- Blue boxes are what **else** you need to build it (about 600)
- E4S sits on top of *hundreds* of open source community projects

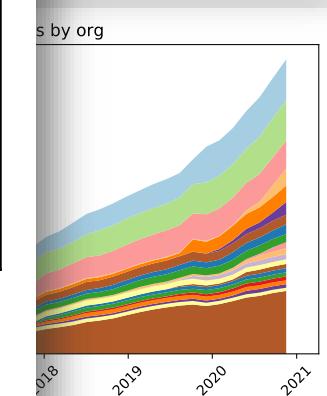
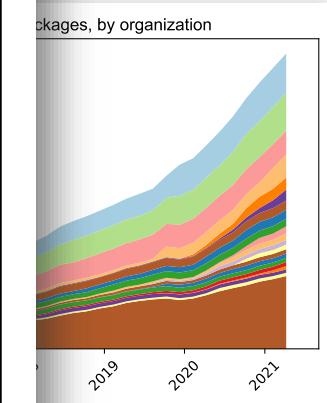
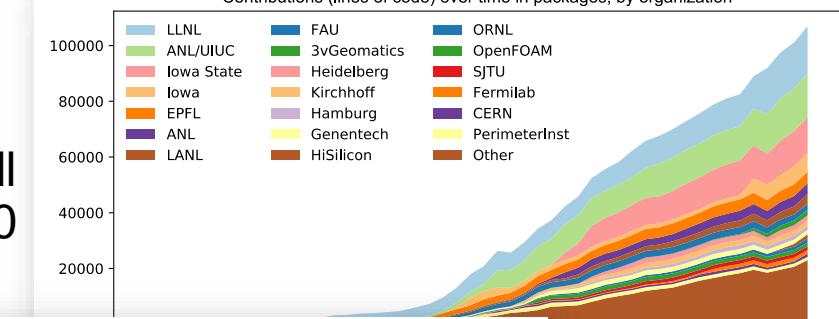
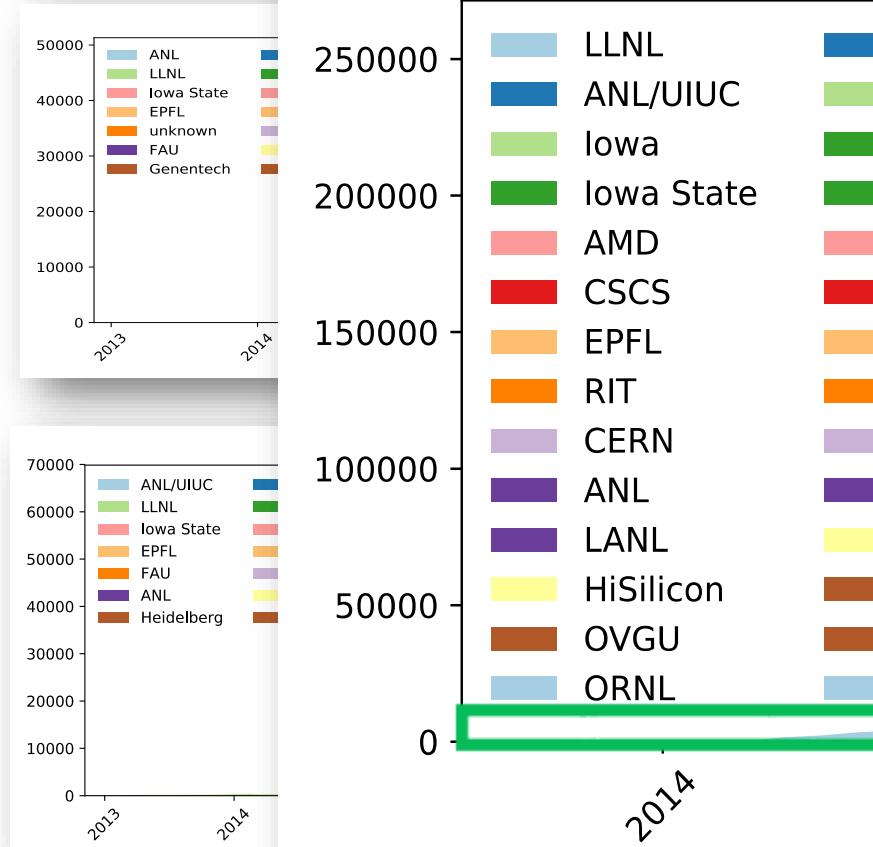


2016

Fall
2020

2024

Contributions (lines of code) over time in packages, by organization



Extreme-scale Scientific Software Stack (E4S)

<https://spack.io>

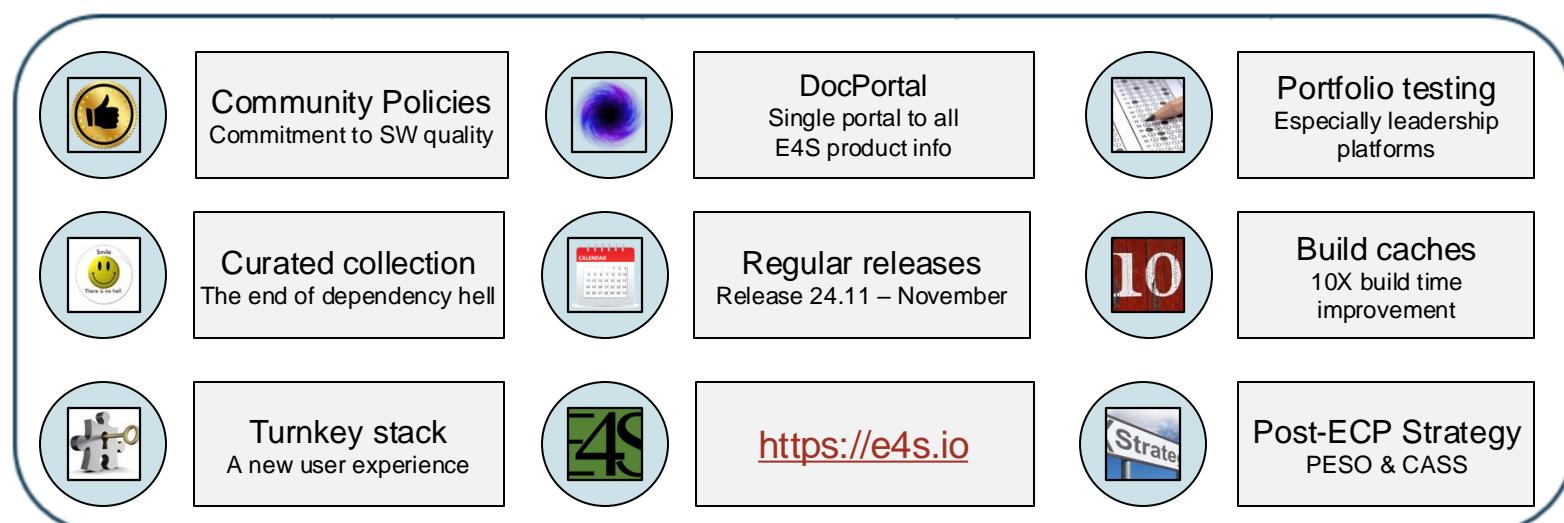
Spack lead: Todd Gamblin (LLNL)

- E4S: HPC software ecosystem – a curated software portfolio
- A **Spack-based** distribution of software tested for interoperability and portability to multiple architectures
- Available from **source, containers, cloud, binary caches**
- Not a commercial product – an open resource for all
- Growing functionality: May 2024: E4S 24.05 – **140+ full release products**
- **One of the most important legacies of ECP**



<https://e4s.io>

E4S lead: Sameer Shende (U Oregon)



Also includes other products, e.g.,
AI: PyTorch, TensorFlow, Horovod
Co-Design: AMReX, Cabana, MFEM

Fast Follower Case Study: DoD Use of E4S

Winter 2024 ECP Industry and Agency Council Meeting

31 January 2024

Highlights from Sean Ziegeler talk at ECP IAC Meeting

- Port and use of Spack/E4S without DOE assistance
- Phase 1:
 - Five HPC systems
 - 3-4 compiling systems
 - 2 MPIs
 - 100 products + dependencies
- Phase 2:
 - Deep dive with 16 products
 - Used by DoD, in E4S

Fast Follower Case Study: DoD Use of E4S

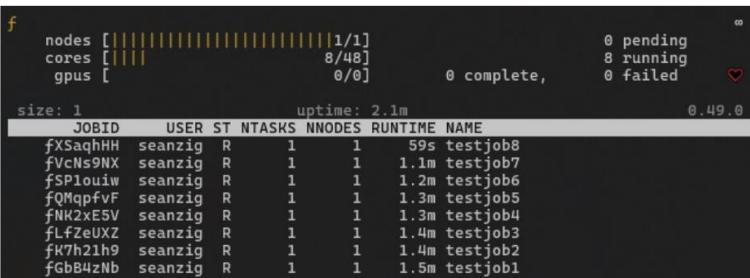
Winter 2024 ECP Industry and Agency Council Meeting

31 January 2024

Selected Results: Flux

- A scheduler (like PBS) but can be launched by the user and manage a set of nodes assigned to the user through the actual system scheduler
 - Not currently installed on HPCM systems
- Install


```
spack install flux-core %gcc
```
- Exciting option for users with lots of miscellaneous tasks to run



The screenshot shows the Flux command-line interface. At the top, it displays resource usage: nodes [1/1], cores [8/48], and gpus [0/0]. It also shows statistics: 0 pending, 8 running, 0 complete, and 0 failed. Below this, it shows the size of the system as 1, uptime as 2.1m, and a timestamp of 0.49.0. The main part of the screen is a table listing jobs:

JOBID	USER	ST	NTASKS	NNODES	RUNTIME	NAME
fXSaqhhH	seanzig	R	1	1	59s	testjob8
fVcNs9NX	seanzig	R	1	1	1.1m	testjob7
fSP1ouiw	seanzig	R	1	1	1.2m	testjob6
fQMqpfvF	seanzig	R	1	1	1.3m	testjob5
fNK2xe5V	seanzig	R	1	1	1.3m	testjob4
fLFzeUXZ	seanzig	R	1	1	1.4m	testjob3
fK7h21h9	seanzig	R	1	1	1.4m	testjob2
fGbB4zNb	seanzig	R	1	1	1.5m	testjob1

Highlights from Sean Ziegeler talk at ECP IAC Meeting

- Port and use of Spack/E4S without DOE assistance
- Phase 1:
 - Five HPC systems
 - 3-4 compiling systems
 - 2 MPIs
 - 100 products + dependencies
- Phase 2:
 - Deep dive with 16 products
 - Used by DoD, in E4S

Fast Follower Case Study: DoD Use of E4S

Winter 2024 ECP Industry and Agency Council Meeting

31 January 2024

Selected Results: Flux



Selected Results: TAU

- **Performance analysis tool**
 - Installed on HPCMP systems as part of Computational Science Environment (CSE)
- **Install options:** `+papi +fortran +io +pdt +pthreads +mpi +openmp +craycnl`
- **Potentially useful options:** `+comm +python +shmem`
- **For GPUs:** `+cuda +rocm +rocprofiler +roctracer`
- **Easier to build and configure versus CSE TAU installation:**
 - Configuring for CUDA works and is straightforward
 - Spack can be configured to use a newer PAPI library which has a wider variety of counters
 - Performance measurements from Spack & CSE versions were comparable
 - Tau_exec from Spack installation was more reliable
 - Except: had to modify the PAPI package script to include options not offered by default

Highlights from Sean Ziegeler talk at ECP IAC Meeting

- Port and use of Spack/E4S without DOE assistance
- Phase 1:
 - Five HPC systems
 - 3-4 compiling systems
 - 2 MPIs
 - 100 products + dependencies
- Phase 2:
 - Deep dive with 16 products
 - Used by DoD, in E4S

Fast Follower Case Study: DoD Use of E4S

Winter 2024 ECP Industry and Agency Council Meeting

31 January 2024

Selected Results: Flux



Selected Results: TAU

Selected Results: Kokkos & RAJA

- C++ performance-portability frameworks
 - Not provided on HPCMP systems
 - Exception is Kokkos as part of LAMMPS

Install

```
spack install --keep-stage kokkos %gcc@12.1.0 +cuda cuda_arch=70
               +openmp +examples +wrapper +tests    # Cray EX
spack install --keep-stage raja %gcc@12.1.0 +cuda cuda_arch=70
               +openmp +examples +exercises +tests
```

Both built and worked effectively on the test systems

- Kokkos-kernels used to test: GPU (V100) runs 2.6x-7.8x faster than CPU (Dual AMD EPYC 7H12, 128 cores)
- Built-in tests: GPU runs 2.3-50x faster than CPU

Highlights from Sean Ziegeler talk at ECP IAC Meeting

- Port and use of Spack/E4S without DOE assistance
- Phase 1:
 - Five HPC systems
 - 3-4 compiling systems
 - 2 MPIs
 - 100 products + dependencies
- Phase 2:
 - Deep dive with 16 products
 - Used by DoD, in E4S

Fast Follower Case Study: DoD Use of E4S

Winter 2024 ECP Industry and Agency Council Meeting

31 January 2024

Selected Results: Flux



Selected Results: TAU



Selected Results: Kokkos & RAJA



Selected Results: LAMMPS



- Classical molecular dynamics code with a focus on materials modeling

- Provided on some systems as needed, most frequently installed manually by users due to customization
 - A frequent customization is represented in the spack install flags ([green](#)) here
 - Enterprise version has the following LAMMPS options also turned on: colloid, kspace, manybody, molecule, phonon, rigid, and shock, and is an MPI build

- Install

```
spack install lammps %oneapi@2022.2.0 -ffmpeg -jpeg -kim -png    # Cray EX
spack install lammps %aocc@3.0.0 -ffmpeg -jpeg -kim -png        # Cray EX
spack install lammps %intel@19.0.4.243 -ffmpeg -jpeg -kim -png # HPE SGI
```

- Spack build is 2X faster than HPCM build for initial tests

- Cause of the performance increase is currently unknown, but some possibilities have been ruled out

Highlights from Sean Ziegeler talk at ECP IAC Meeting

- Port and use of Spack/E4S without DOE assistance
- Phase 1:
 - Five HPC systems
 - 3-4 compiling systems
 - 2 MPIs
 - 100 products + dependencies
- Phase 2:
 - Deep dive with 16 products
 - Used by DoD, in E4S

Fast Follower Case Study: DoD Use of E4S

Winter 2024 ECP Industry and Agency Council Meeting

31 January 2024

Selected Results: Flux



Selected Results: TAU



Selected Results: Kokkos & RAJA



Selected Results: LAMMPS



- Classical molecular dynamics code with a focus on materials modeling

- Provided on some systems as needed, most frequently installed manually by users due to customization
 - A frequent customization is represented in the spack install flags ([green](#)) here
 - Enterprise version has the following LAMMPS options also turned on: colloid, kspace, manybody, molecule, phonon, rigid, and shock, and is an MPI build

- Install

```
spack install lammps %oneapi@2022.2.0 -ffmpeg -jpeg -kim -png    # Cray EX
spack install lammps %aocc@3.0.0 -ffmpeg -jpeg -kim -png        # Cray EX
spack install lammps %intel@19.0.4.243 -ffmpeg -jpeg -kim -png # HPE SGI
```

- Spack build is 2X faster than HPCMP build for initial tests

- Cause of the performance increase is currently unknown, but some possibilities have been ruled out

Highlights from Sean Ziegeler talk at ECP IAC Meeting

- Port and use of Spack/E4S without DOE assistance
- Phase 1:
 - Five HPC systems
 - 3-4 compiling systems
 - 2 MPIs
 - 100 products + dependencies
- Phase 2:
 - Deep dive with 16 products
 - Used by DoD, in E4S

E4S 24.05 Release

- **140+ HPC-AI packages** on ARM, x86_64, ppc64le platforms, 128K+ binaries in E4S build Cache
- **Growing suite of AI/ML packages**
 - DeepHyper, Google.generativeai (Gemini API), OpenAI (API), TorchBraid, Pandas, Scikit-Learn, JAX, PyTorch, TensorFlow, Horovod, OpenCV, and LBANN with support for GPUs
 - E4S DocPortal updated with AI/ML tools
- OS upgrade for containers: **Ubuntu 22.04 LTS**
- Upgraded
 - **CUDA from version 11 to 12,**
 - **ROCM upgraded from version 5.4 to 5.7.1.**
- **New tools:** Laghos, Glvis, netcdf-fortran, fpm, e4s-cl, and e4s-alc.
- **New applications:** Nek5000, Nekbone, Laghos and previously supported GROMACS, CP2K, Xyce, Quantum Espresso, ExaGo, LAMMPS, WARPX, Dealii, and OpenFOAM.
- **Adaptive Computing's HPC Cloud on-demand data center (ODDC) web-based platform**
 - Multi-user, multi-node ParaTools Pro for E4S images on AWS marketplace with support for aarch64 (Graviton) as well as x86_64 with NVIDIA GPUs with
 - VNC based remote desktop and torque (qsub) for multi-node execution
 - <https://adaptivecomputing.com/>

CI Testing:

- Frank: Extensive multi-device system
- Includes Grace-Grace, Grace-Hopper
- 7M+ CI builds

E4S 24.11 Release – Look for SC24 announcement

Looking Forward Post-ECP Software Stewardship and Advancement

8 Software Stewardship Organizations (SSOs)

DOE Office of Advanced Scientific Computing Research (ASCR) Post-ECP Projects

COLABS

Training, workforce development, and building the RSE community

CORSA

Partnering with foundations to provide sustainable pathways for scientific software

FASTMATH

Stewardship, advancement, and integration for math and ML/AI packages

PESO

Stewarding, evolving and integrating a cohesive ecosystem for DOE software

RAPIDS

Stewardship, advancement, and integration for data, visualization and ML/AI packages

S4PST

Stewardship, advancement and engagement for programming systems

STEP

Stewardship, advancement of software tools for understanding performance and behavior

SWAS

Stewardship and project support for scientific workflow software and its community



Announcing CASS

The Consortium for the Advancement of Scientific Software
<https://cass.community>

CASS Basics

- A newly-formed organization
- Sponsored by DOE Office of Advanced Scientific Computing Research (ASCR)
- Established by DOE Software Stewardship Organizations (SSOs)

CASS Goals

- Forum for SSO collaboration and coordination
- Bigger than the sum of its parts
- Vehicle for advancing the scientific software ecosystem

CASS Status

- Defining governance structure
- Establishing community awareness
- Building a team of teams
- Collaborating on outreach

Software Stewardship Organization (SSO) Basics

- Each SSO represents a specific software ecosystem concern
- **Product SSOs:** Programming systems, performance tools, math packages, data/viz packages
- **Portfolio SSO:** Curating & delivering software stack to the community
- **Community SSOs:** Workforce, partnerships

Engage with CASS

- Review slides June 11-13 CASS Community BOF Days:
<https://cass.community/bufs>
- Visit <https://cass.community>

Consortium for the Advancement of Scientific Software (CASS)

- **CASS Formation Team**
 - David Bernholdt
 - Phil Carns
 - Lois Curfman McInnes
- **Representatives of Software Stewardship Organizations (SSOs)**
 - **COLABS:** Anshu Dubey, David Bernholdt
 - **CORSA:** Greg Watson, Elaine Raybourn
 - **FASTMath:** Esmond Ng, Todd Munson
 - **PESO:** **Mike Heroux, Todd Gamblin**
 - **RAPIDS:** Rob Ross, Lenny Oliker
 - **S4PST:** Keita Teranishi, Damian Rouson
 - **STEP:** Terry Jones, Phil Carns
 - **SWAS:** Rafael Ferreira da Silva, Lavanya Ramakrishnan

PESO: Partnering for Scientific Software Ecosystem Stewardship Opportunities

About PESO

- Five-year post-ECP software-ecosystem stewardship and advancement project
- Partnership with CASS
- PESO leads portfolio activities of E4S+Spack curation, testing, delivery

Key PESO goals

- Enable applications to realize benefits of a software ecosystem
- Emphasize software product quality, the continued fostering of software product communities, and the delivery of products, working with CASS

Key PESO Activities

- **Partnerships:** We lead CASS efforts for diverse and inclusive workforce with sustainable career paths. We shepherd BSSw Fellows Program and BSSw.io portal.
- **Services:** We provide services including software product management, integration, and delivery, as well as software quality assurance and security.
- **Products:** We deliver & support products via Spack & E4S, provide porting & testing platforms leveraged across product teams to ensure code stability & portability.

PI: Michael Heroux, Sandia National Laboratories

Collaborating Institutions: ANL, Berkeley, BNL, Kitware, LLNL, LANL, ORNL, PNNL, SNL, SHI, UO

ASCR Program: Software Stewardship and Advancement

ASCR PM: William Spotz

Resources: <https://pesoproject.org>, <https://e4s.io>, <https://hpsf.io>



The PESO Project exists to preserve, sustain, and advance the investments made by the Exascale Computing Project in a robust, versatile, and portable HPC software ecosystem and the people who make the ecosystem effective.

PESO is unique in its organization by composing itself of many members of other software stewardship organizations (SSOs) to best ensure tight integration across the SSOs and support the mission of the Consortium for the Advancement of Scientific Software (CASS).

Are these challenges familiar? Investing in the Spack/E4S ecosystem can help

- Installing 3rd party libs and tools
- Managing library version updates
- Adopting new libraries and tools
- Managing portability across CPUs/GPUs
- Long build times
- Out-of-date algorithms
- Awareness of latest SW practices
- Opportunities to improve your software

**Spack and E4S
“Powers of 10”
possibilities**

100,000+

Lines of code replaced with high-quality libs & tools

Using robust open-source libraries and tools helps you eliminate native source code that is less capable, more fragile, and harder to maintain

100+

Speedup using advanced devices like GPUs

ECP-sponsored applications realize a factor of 100 or more science-speedup by reformulating algorithms and software to exploit GPUs

10,000+

Community members via ecosystem collaborations

Shared experiences using the same ecosystem increases community knowledge base, improved software capabilities, cross-teaming

10+

Reduction in build times via Spack build caches

Spack build caches reduce re-build times by a factor of 10 or more, greatly reducing staff wait times and accelerating debugging and reconfiguration activities

1,000+

Code teams share ecosystem costs & benefits

Pooling investments and making the software ecosystem available everywhere optimizes cost and benefit sharing across many users

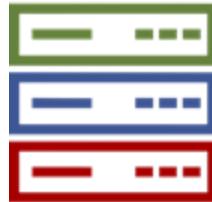
1

Source code base for all computing systems

Your single source code base that achieves performance portability, now and in the future, using libraries, tools, and compilers available via Spack and E4S

The High-Performance Software Foundation

We launched the High Performance Software Foundation (HPSF) at ISC24



HIGH PERFORMANCE SOFTWARE FOUNDATION

Scientific Achievements

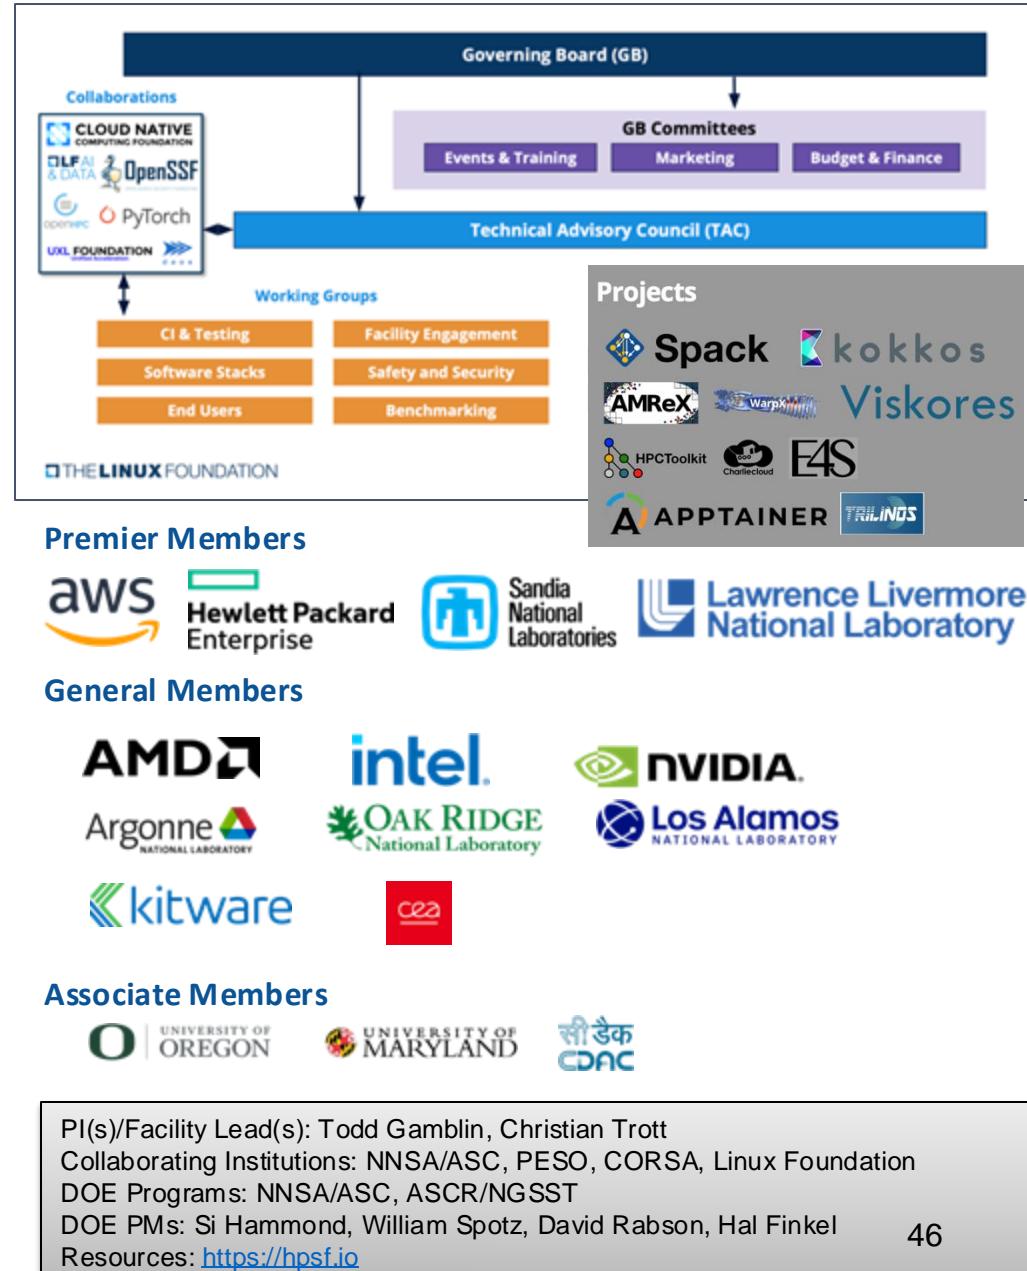
- NNSA led the formation of HPSF over the past 18 months, in close collaboration with DOE/ASCR, industry, and the Linux Foundation
- 15 founding members, 6 initial projects from industry, academia, labs around the world
- HPSF provides open source projects with:
 - A neutral home
 - Collaborations
 - Open governance
 - Funds for project infrastructure, working groups, events, other initiatives

Significance and Impact

- Initial membership raised more funds than expected: **~\$1M annual budget**
- ISC kickoff BOF session was standing-room only; generating much excitement
- Expect to grow membership and projects over time
 - 2 NNSA, 2 ASCR, 2 European projects in the pipeline
 - Expecting at least 2 more general members within the year

Approach

- Separate financial (GB), technical (TAC), and project governance
- Interact directly with key projects in the HPC ecosystem
- Grow contributor base through increased adoption, training, events, outreach
- Build a portable, accelerated software stack for HPC and beyond



Some opportunities going forward

- The growing success of generative AI is spurring disruptive changes in scientific software
- GPUs are really AI devices (CPUs are still relevant and can benefit from concurrency too)
- A recent Hyperion survey estimates the high-end AI market is **30 times larger** than traditional HPC
 - Bad news: We are no longer a strong market driver
 - Good news: The HPC market is still large but the AI market is huge. We can benefit greatly, if we want ...
- Here are three ways to benefit:
 1. **New algorithms adapted to AI devices**
New computing devices, such as GPUs designed for AI workflows, are suitable for traditional scientific codes but require extensive algorithm and software changes to realize their potential.
 2. **AI models to complement or replace analytic models**
AI inference engines are compelling components, complementing or replacing traditional modeling and simulation approaches.
 3. **Leverage AI tools in our research work**
Generative AI tools are transforming the entire research enterprise, especially software activities. AI tools assist in producing requirements, specifications, designs, source code, tests, and more. These changes are exciting but come with risks.

Summary

- The Exascale Computing Project (ECP):
 - Explored many approaches to performance-portable accelerator-based computing
 - Focused on very high end but most innovations were on the node: suitable for desktop on up
 - De-risked the path for many others: 62 applications from many domains showed a path forward
- In addition to transforming your own software base, you could consider using ECP-developed products:
 - Kokkos, zfp, Spack are a few
 - E4S provides many others
 - HPSF-sponsored products are also emerging
- PESO, CASS, and other software stewardship projects are evolving to meet future needs
- Happy to talk further: mheroux@acm.org

Post-ECP and Final Remarks

DOE/ECP has **learned a lot about producing software contributions** to the HPC community:

- Improved planning, executing, tracking, assessing, integrating, and delivering
- Improved interactions with the broader HPC software and hardware community
- Direct engagement with industry, US agencies, and international collaborators

In post-ECP efforts **we propose to continue and expand these efforts:**

- Further engage with commercial partners to provide a rich, robust software ecosystem
- Evolve a stable, sustainable business model for engaging with agencies and industry
- Engage with cloud providers, software foundations, and others to optimize cost & benefit sharing
- Further the ECP strategy for direct industry and agency engagement

We intend to **realize the potential of the ECP legacy across the HPC community:**

- **Realize the potential of software ecosystems** by leveraging “powers of 10” advantages
- **Increase the trustworthiness, sustainability, and cost effectiveness** of our software in the future
- **Support existing and emerging needs for AI/ML libraries and tools**

We want to work with the HPC community to realize the legacy of ECP, and beyond

- We have many new ways to interact
- Many new opportunities to pursue

Thank you!